



Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

Lecture 1 Introduction

Shiwei Lan¹

¹School of Mathematical and Statistical Sciences
Arizona State University

STP598 Advanced Deep Learning Models
Spring 2026



Course Info

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

Advanced Deep Learning Models

- Instructor: Shiwei Lan <[slan@asu.edu](mailto:silan@asu.edu)>
 - Office: WXLR 544
 - Office hours: TuTh 12 - 1:15p @ <https://asu.zoom.us/j/8055899886>
- Teaching Assistant:
 - Ping-Han Huang <phuang37@asu.edu>
 - Zoom: <https://asu.zoom.us/j/87568448130?jst=2>
 - Office hours: TBA



Course Info

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- Course Schedule, Lecture Notes, etc.

<https://slan-teaching.github.io/STP598adl/>

- Discussion and Questions

Canvas Discussions [STP 598adl](#)

- Homework and Grades on [Canvas](#)

<https://canvas.asu.edu/courses/248838>

- Coding Assignments on [Nbgrader](#)

https://agnesi.la.asu.edu/services/stp598_28057



Table of Contents

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

① Deep Learning: Opportunities and Challenges

② Course Overview

③ DNN Review

④ Computing

What is Deep Learning?

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

Deep learning is a branch of machine learning that uses algorithms called **neural networks**—inspired by the human brain—to learn patterns from large amounts of data.





What is Deep Learning?

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- A neural network is made of layers of connected “neurons” (mathematical functions).
- Deep learning means the network has many layers (input → multiple hidden layers → output).
- Each layer learns increasingly complex features:
 - In images: edges → shapes → objects
 - In language: letters → words → meaning
- The model learns by adjusting millions (or billions) of parameters to minimize errors.

Deep learning automatically *learns useful features from raw data*, instead of relying on humans to hand-engineer them.

What does deep learning do?

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- ① It handles complex, high-dimensional data
 - It works extremely well with images and video, Speech and audio, natural language (text, chatGPT), and sensor data (self-driving, robotics).
- ② It reduces manual feature engineering
 - Before: experts manually decide what features matter; After: the model learns the features by itself.
- ③ It improves accuracy at scale
 - When you have a lot of data and computing power, deep learning usually outperforms other methods. (Scaling law).
- ④ It enables modern AI applications
 - Image recognition (face unlock, medical imaging), voice assistants (Google, Alexa), FSD (Tesla), LLM (ChatGPT, Gemini, Claude).

What challenges does deep learning face?

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

① Vanishing and Exploding Gradients

- As networks became deeper, gradients during backpropagation would either vanish or explode. Solution: activation functions, initialization schemes, normalization, residual learning.

② Overfitting and Poor Generalization

- Large networks memorize training data and perform poorly on unseen data. Solution: regularization (dropout), data augmentation, pre-training+fine-tuning (transfer-learning), self-supervised learning.

③ Data Hunger and Label Dependence

- Self-supervised and contrastive learning (BERT-style masked modeling), foundation models.

④ Limited Long-Range Dependency Modeling

- CNNs, RNNs or LSTMs struggle with long sequences and global context. Solution: transformers and attention mechanisms, sparse and linear attention variants.



What challenges does deep learning face?

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

5 Training Instability and Optimization Difficulty

- Training is sensitive to learning rates, initialization, hyperparameters. Solution: advanced optimizers, learning rate schedules, and gradient clipping.

6 Computational Inefficiency

- Hardware acceleration (GPUs, TPUs, NPUs), parallelism, efficient architectures (MoE).

7 Lack of Interpretability

- DNNs are "black boxes" and hard to understand or trust. Solution: attention visualization, mechanistic interpretability, and physics informed learning.

8 Task-Specific, Narrow Models

- Trained for a single task, hard to reuse. Solution: foundation models, multitask and multimodal learning, and prompt-based adaptation.



Top AI Achievements: what are your votes?

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- ① Generative AI (e.g., GPT, DALL-E, Stable Diffusion)
 - generative AI: Generative artificial intelligence (Wikipedia)
- ② ChatGPT and Advanced Language Models
 - AI Breakthrough Timeline
- ③ AlphaGo and Game-Playing AI
 - 70 Years of AI: Key Milestones
- ④ AlphaFold — Protein Structure Prediction
 - AlphaFold Changed Science
- ⑤ Advances in Multimodal and Reasoning AI
 - AI Breakthrough Timeline



Table of Contents

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

1 Deep Learning: Opportunities and Challenges

2 Course Overview

3 DNN Review

4 Computing



Course Overview

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- **Residual Networks (ResNet)**: enabling training of very deep NNs
- **U-Net**: fundamental CNN architecture for image segmentation
- **General Adversarial Networks (GAN)**: introducing adversarial training to generate realistic images
- **Transformer**: introducing self-attention mechanism to enable parallelism and scaling, backbone of modern foundation models
- **Bidirectional Encoder Representations from Transformers (BERT)**: shifting NLP from task-specific models to more general
- **Diffusion Models**: dominant generative models
- **Graph Neural Networks (GNN)**: network for graph-structured data
- **Large language models (LLMs)**: modern AI applications



Table of Contents

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

1 Deep Learning: Opportunities and Challenges

2 Course Overview

3 DNN Review

4 Computing

Overview

S.Lan

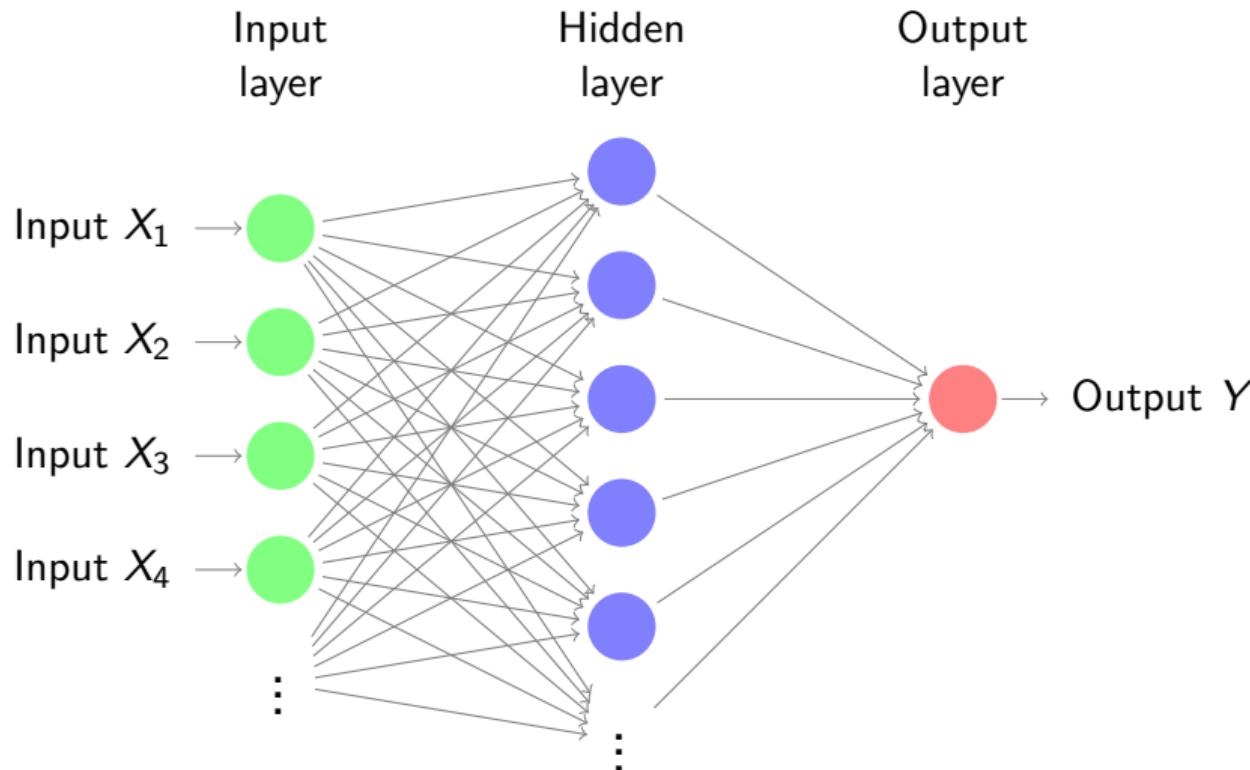
Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- Neural Networks were first developed as models for the human brain, where we have many units (**neurons**) that simultaneously process signals to give a joint decision.
- The neurons fire when the total signal passed to that unit exceeds a certain threshold.
- The collective signal from all neurons tells you whether its a dog or a cat.

[Overview](#)[S.Lan](#)[Deep Learning:
Opportunities
and Challenges](#)[Course Overview](#)[DNN Review](#)[Computing](#)



Feedforward neural network

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- Given a training set $\{x_i, y_i\}_{i=1}^n$,
 - For regression: $y_i \in \mathbb{R}^K$ is a K dimensional continuous outcome
 - For classification: $y_i \in \{1, 2, \dots, K\}$
- The goal is still to model the relationship

$$E(Y|X) = f(X)$$

- Instead of modeling the probabilities directly using X , we build **M hidden neurons** as a hidden layer between X and Y :

$$\begin{aligned}Z &= (1, Z_1, Z_2, \dots, Z_M) \\&= (1, \sigma(\mathbf{X}^\top \boldsymbol{\alpha}_1), \sigma(\mathbf{X}^\top \boldsymbol{\alpha}_2), \dots, \sigma(\mathbf{X}^\top \boldsymbol{\alpha}_M))\end{aligned}$$

[Overview](#)[S.Lan](#)[Deep Learning:
Opportunities
and Challenges](#)[Course Overview](#)[DNN Review](#)[Computing](#)

- $\sigma(\cdot)$ is an **activation function**. Some examples?
- We model Y using the hidden layer variables Z through some **link function** $g(\cdot)$

$$X \xrightarrow{\sigma(\cdot)} Z \xrightarrow{g(\cdot)} Y$$

- In **classification problems** (K class), we can use logit link g_k to model the probability of $Y = k$, for $k = 1, \dots, K$:

$$g_k(Z) = \frac{\exp(Z^T \beta_k)}{\sum_{l=1}^K \exp(Z^T \beta_l)}$$

- In **regression problems** (could be multidimensional), we can simply use a linear function to model the k th entry of Y :

$$g_k(Z) = Z^T \beta_k$$



Feedforward neural network

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- The multidimensional function $\mathbf{f}(x)$ can be represented as a convoluted way of mapping $x \in \mathbb{R}^P$ to $y \in \mathbb{R}^K$

$$\mathbf{f}(x) = \mathbf{g} \circ \sigma(x)$$

- The notations \mathbf{g} and σ here are multidimensional.
- The parameters involved are: $\alpha_1, \dots, \alpha_M$, and β_1, \dots, β_K .



Examples of activation functions

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- The **activation function** $\sigma(\cdot)$ takes a linear combination of the input variables, and output a scalar through nonlinear transformation. Examples:

- sigmoid:

$$\sigma(v) = \frac{1}{1 + e^{-v}} = \frac{e^v}{e^v + 1}$$

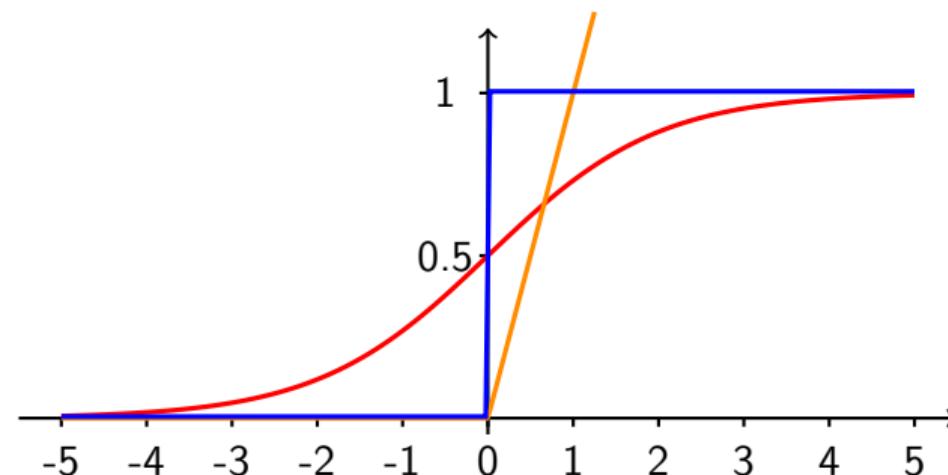
- hyperbolic tangent (tanh):

$$\sigma(v) = \frac{e^v - e^{-v}}{e^v + e^{-v}}$$

- rectified linear unit (ReLU):

$$\sigma(v) = \max(0, v), \quad \text{soft approx. } \ln(1 + e^v)$$

- And many others: exponential linear unit, arctangent, etc.

[Overview](#)[S.Lan](#)[Deep Learning:
Opportunities
and Challenges](#)[Course Overview](#)[DNN Review](#)[Computing](#)

Sigmoid: $(1 + e^{-v})^{-1}$

ReLU: $\max(0, v)$,

Step function: $I(v > 0)$



Activation Functions

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- Originally, a **step function** $I(v > 0)$ was considered as the activation function (to mimic the biological interpretation). Hence for each neuron, signal is triggered only when $x^T \alpha$ is above a certain threshold
- It was later recognized that the step function is **not smooth** enough for optimization, hence was replaced by a smoother threshold function, the sigmoid function
- “Feedforward” as signals can only pass to the next layer. There is no “cycle” in the model



Why Neural Networks work

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

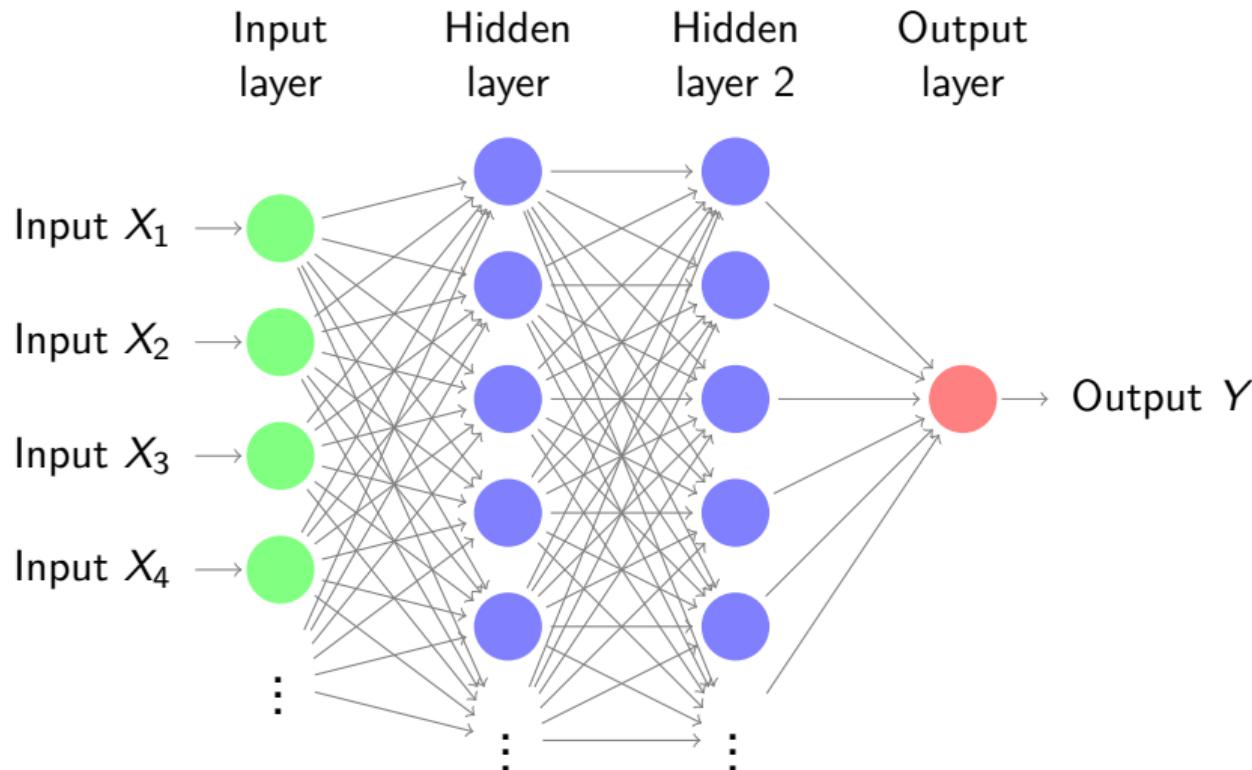
Computing

Universal Approximation Theorem (Cybenko, 1989; Hornik 1991)

Any continuous function $f(x)$ on the space $[0, 1]^P$ can be approximated (for any $\epsilon > 0$) by a finite set of neurons with a bounded monotone-increasing activation function $\sigma(\cdot)$:

$$|f(x) - \sum_k w_k \sigma(\beta_k^T x + b_k)| < \epsilon$$

for some w_k , β_k , and b_k . Hence, the functions defined by the neurons is dense.

[Overview](#)[S.Lan](#)[Deep Learning:
Opportunities
and Challenges](#)[Course Overview](#)[DNN Review](#)[Computing](#)



Neural Networks

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- Try this a really cool website: <http://playground.tensorflow.org/>
- Implementation in Python :
 - packages: `sklearn.neural_network`, (TensorFlow, PyTorch)
 - `MLPClassifier` implements a multi-layer perceptron (MLP) algorithm for classification.
 - `MLPRegressor` implements a multi-layer perceptron (MLP) for regression.
 - MLP trains using Stochastic Gradient Descent , Adam , or L-BFGS .
 - Important parameters:
 - number of neurons: `hidden_layer_sizes`
 - activation functions: `activation`
 - size of minibatches: `batch_size`
 - solver for back-propagation: `solver`
 - learning step sizes: `learning_rate`, `learning_rate_init`
 - regularization: `alpha`



Fitting Neural Networks

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- The parameters (weights) α 's and β 's need to be optimized.
- For a single hidden layer NN, we have

$$\{\alpha_1, \dots, \alpha_M\} : M(p+1) \text{ weights}$$

$$\{\beta_1, \dots, \beta_K\} : K(M+1) \text{ weights}$$

- where p is the number of non-intercept X features; M is the number of hidden neurons in a single layer; and K is the number of categories for classification.
- $K = 1$ if its a univariate regression problem.



Fitting Neural Networks

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- Neural Networks training is based on error minimization using a **Gradient Descent** algorithm, known as error **back-propagation**.
- For K classification, we minimize Deviance:

$$-\sum_{i=1}^n \sum_k \mathbf{1}\{y_i = k\} \log f_k(x_i)$$

- For univariate regression, we minimize RSS (since g is linear):

$$\sum_{i=1}^n (y_i - f(x_i))^2 = \sum_{i=1}^n (y_i - \beta_0 - \beta_1 \sigma(x^\top \alpha_1) - \cdots - \sigma(x^\top \alpha_M))^2$$



Fitting Neural Networks

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- The objective function can be written as

$$R(\theta) = \sum_{i=1}^n R_i(\theta)$$

where R_i represents the deviance or residual sum of squares for the i th data point, and θ represents an aggregated vector of all weights

- Initiate weights $\theta^{(0)}$
- We then calculate the derivative wrt each of the weights evaluated at the current iteration value $\theta^{(t)}$:

$$\sum_{i=1}^n \frac{\partial R_i}{\beta_{km}} \Big|_{\theta=\theta^{(t)}} \quad \sum_{i=1}^n \frac{\partial R_i}{\alpha_{mj}} \Big|_{\theta=\theta^{(t)}}$$

- Stochastic GD:** the summation can be taken over a **random subset** of the n samples

GD vs. Stochastic GD

Overview

S.Lan

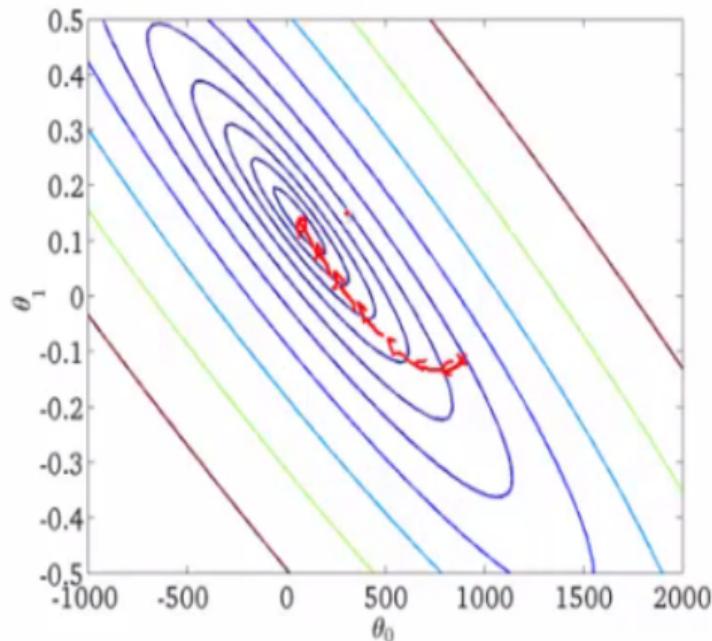
Deep Learning:
Opportunities
and Challenges

Course Overview

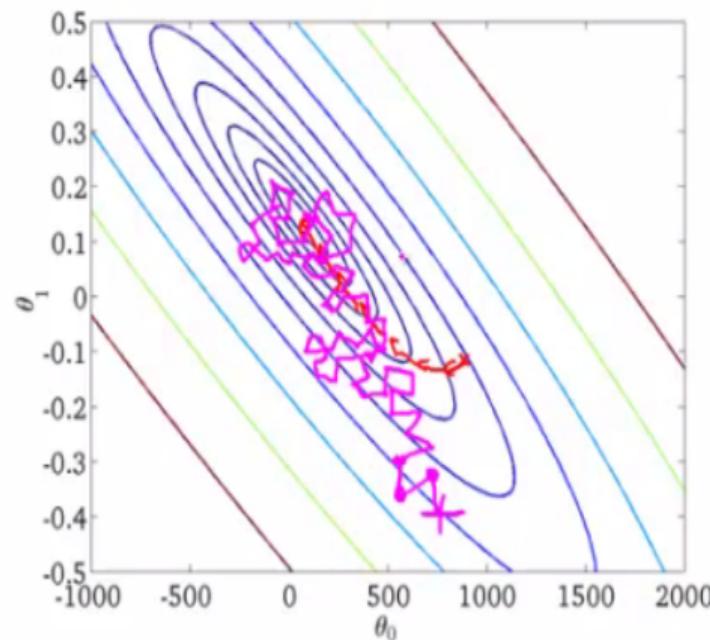
DNN Review

Computing

Gradient Descent vs.



Stochastic Gradient Descent





Fitting Neural Networks

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- The derivatives for $K = 1$ regression case is essentially

$$\frac{\partial R_i}{\beta_m} = -2(y_i - f(x_i))z_{mi}$$

$$\frac{\partial R_i}{\alpha_{ml}} = -2(y_i - f(x_i))\beta_m \sigma'(\alpha_m^T x_i) x_{il}$$

- Some redundant calculations can be saved in the above equations. The property is called **back-propagation**.
- We then do the update, at the t -th iteration

$$\beta_m^{(t+1)} = \beta_m^{(t)} - \gamma \sum_{i=1}^n \frac{\partial R_i}{\beta_m^{(t)}}$$

$$\alpha_{ml}^{(t+1)} = \alpha_{ml}^{(t)} - \gamma \sum_{i=1}^n \frac{\partial R_i}{\alpha_{ml}^{(t)}}$$

where γ is a step size for gradient descent.



Fitting Neural Networks

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- The derivatives can be calculated by Chain Rules
- The algorithm can be implemented by a forward-backward sweep over the network
- In the **forward** pass, compute the hidden variables and the output $\hat{f}(x_i)$ based on the current weights $\theta^{(t)}$
- In the **backward** pass, compute the derivatives, and update $\theta^{(t)} \rightarrow \theta^{(t+1)}$



Going Deeper...

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- Deep Neural Networks are one type of deep learning models.
- Deep neural Networks are just ... Neural Networks with more than one hidden layer.
- But neural networks have been around for more than 70 years... why it gets popular just in recent years?
 - computational issues
 - a better way to generate/construct features
 - ...

[Overview](#)[S.Lan](#)[Deep Learning:
Opportunities
and Challenges](#)[Course Overview](#)[DNN Review](#)[Computing](#)

Deep neural networks learn hierarchical feature representations

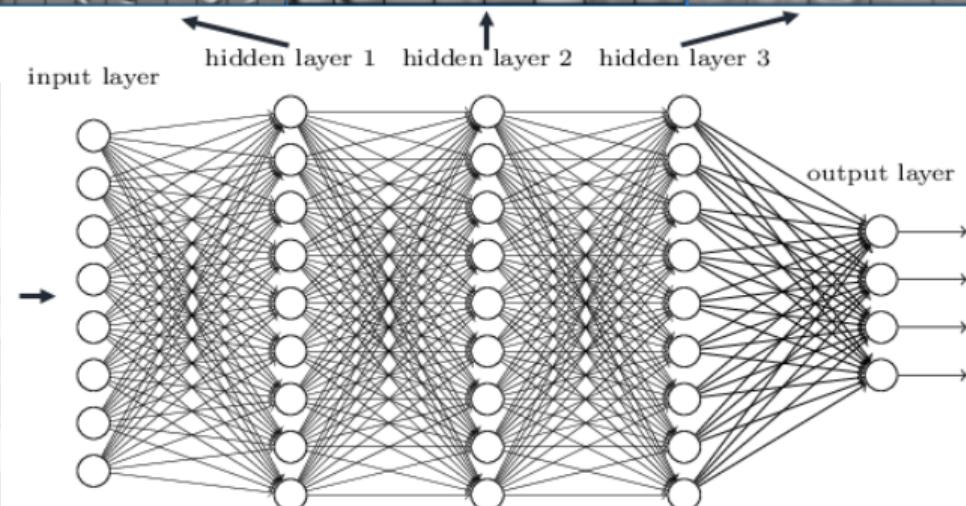
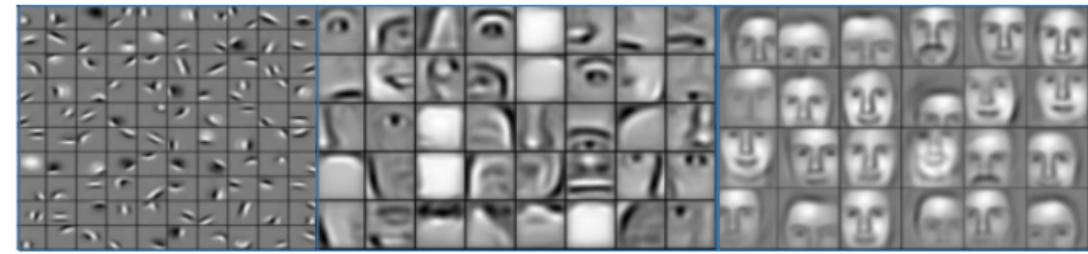




Table of Contents

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

1 Deep Learning: Opportunities and Challenges

2 Course Overview

3 DNN Review

4 Computing



Python, Conda and Jupyter Notebook

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- Python is a free software environment for statistical machine learning and graphics.
 - <https://www.python.org>
- Conda is an open source package / environment management system for Python.
 - <https://anaconda.org>
- Jupyter Notebook is a Python package to integrate code, equations, visualizations in document.
 - <https://jupyter.org>
- Coding IDE Both Eclipse (+PyDev) and Visual Studio Code are recommended.
 - <https://www.eclipse.org> (+ <https://www.pydev.org>) or <https://code.visualstudio.com>.
- R Markdown has also integrated Python. Check [link1](#) [link2](#).



ASU Resources

Overview

S.Lan

Deep Learning:
Opportunities
and Challenges

Course Overview

DNN Review

Computing

- ASU Research Computing supports computing resources that can be used for fitting models.
 - <https://cores.research.asu.edu/research-computing>. One can start from [here](#).
- ASU has a dedicated website for Artificial Intelligence.
 - <https://ai.asu.edu>
- ASU provides multiple AI tools
 - <https://ai.asu.edu/ai-tools> You can get chatGPT-5.2 for free!
- Check out the AI resources and Guidelines
 - <https://instruction.thecollege.asu.edu/AIguidelines>